

# Lab-08

*Ehshan Veerabangsa*

*29 November 2017*

## SVM on Linear separable dataset

### Generate Linearly Separable Data

First we will create a dataset of 100 observation evenly split between two classes, positive and negative. Each observation will have 2 features. The mean of the two class will be the only differential.

```
# Set seed
set.seed(300)

# Observations
observations <- 100

# Features
feaures <- 2

# Data points for each class
data_points <- observations * feaures

# Cluster means
pos_mean <- 0
neg_mean <- 3

postive = matrix(rnorm(data_points, mean = pos_mean), nrow = 100, ncol = 2)
negative = matrix(rnorm(data_points, mean = neg_mean), nrow = 100, ncol = 2)
```

We will now simulate labels for the dataset.

```
labels <- c(rep(1, 100), rep(-1, 100))
```

Next, we bind the dataset to one variable and populate a dataframe.


```
obs <- rbind(postive, negative)

data <- data.frame(x = obs, y = as.factor(labels))
```

The next step is to visualise the dataset with a plot.

```
plot(obs, col=ifelse( labels > 0, 1, 2))

legend("topleft",c("Positive","Negative"),col=seq(2),pch=1,text.col=seq(2))
```



Finally we can split our data into training and test sets by a 70/30 split

```
# Split data into training & test sets

set.seed(300)

train_indices <- sample(200, 200 * 0.7)

training_set <- data[train_indices, ]

test_obs <- obs[-train_indices, ]
test_labels <- labels[-train_indices]
```

## Part 1: Load SVM Library

Before we begin our experiments, we will need to load the appropriate library.

```
#install.packages("e1071")
library(e1071)
```

```
## Warning: package 'e1071' was built under R version 3.3.3
```

## Part 2: Linear SVM with cost = 1

Now we can fit a linear SVM with a cost parameter of 1.

```
linear_fit_1 = svm(y~.,data = training_set, kernel = "linear",cost = 1)
```

### Part 3: Report Results

The next step is to check the figures from our new model.

```
summary(linear_fit_1)

##
## Call:
## svm(formula = y ~ ., data = training_set, kernel = "linear",
##      cost = 1)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   1
##   gamma:    0.5
##
## Number of Support Vectors:  15
##
##  ( 7 8 )
##
##
## Number of Classes:  2
##
## Levels:
##  -1 1
```


The results shows that there are 12 support vectors

### Part 4: Linear SVM Plot

We will now visualise our model.

```
plot(linear_fit_1, training_set)
```

## SVM classification plot



The results show that there are 3 points close to the margin, 1 on the pink side, 2 on the blue.

## Part 5: Predictions on Linear Model

Finally for this model, we can check the error rate, by making predictions, and comparing them with our test set.

```
# predictions
prediction_linear_1 = predict(linear_fit_1, newdata = test_obs)

# error calculation
mean(prediction_linear_1 != test_labels)

## [1] 0.01666667
```

The model has an error rate of 0.01666667.

## Part 6: Linear Models with cost 0.01 & 100000

We will now repeat parts 2 - 5 for model with the cost parameter of 0.01 and 1e5

a: cost = 0.01

```
# build model
linear_fit_001 = svm(y~., data = training_set, kernel = "linear", cost = 0.01)
```

```

# summary
summary(linear_fit_001)


##
## Call:
## svm(formula = y ~ ., data = training_set, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##      gamma: 0.5
##
## Number of Support Vectors:  94
##
## ( 47 47 )
##
##
## Number of Classes:  2
##
## Levels:
##  -1 1

# 92 support vectors

# plot
plot(linear_fit_001, training_set)

```

## SVM classification plot



```
# similar boundray

# predictions
prediction_linear_001 = predict(linear_fit_001, newdata = test_obs)

# error calculation
mean(prediction_linear_001 != test_labels)
```

```
## [1] 0.01666667
```

The results show that this model has the same error rate.

**b: cost = 1e5**


```
# build model
linear_fit_1e5 = svm(y~.,data = training_set, kernel = "linear",cost = 1e5)

# summary
summary(linear_fit_1e5)
```

```
##
## Call:
## svm(formula = y ~ ., data = training_set, kernel = "linear",
##      cost = 1e+05)
##
##
```

```
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  1e+05
##       gamma: 0.5
##
## Number of Support Vectors: 10
##
## ( 5 5 )
##
## Number of Classes: 2
##
## Levels:
##   -1 1
# 8 support vectors
# plot
plot(linear_fit_1e5, training_set)
```

**SVM classification plot**



```
# Only 2 points close to boundary on pink side
# predictions
prediction_linear_1e5 = predict(linear_fit_1e5, newdata = test_obs)
# error calculation
```

```
mean(prediction_linear_1e5 != test_labels)
```

```
## [1] 0.01666667
```

Again we have the same error rate.

## Part 7: Find Optimal Model

We can the tune function to find the optimal cost parameter.

```
set.seed(1)

# range of cost values
costs <- c(0.001,0.01,0.1,1,5,10,100,1000,10000,1e5)

tuned_fit = tune(svm, y~., data = data ,kernel="linear",ranges=list(costs))

summary(tuned_fit)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   Var1
##   0.001
##
## - best performance: 0.03
##
## - Detailed performance results:
##   Var1 error dispersion
## 1  1e-03  0.03  0.0421637
## 2  1e-02  0.03  0.0421637
## 3  1e-01  0.03  0.0421637
## 4  1e+00  0.03  0.0421637
## 5  5e+00  0.03  0.0421637
## 6  1e+01  0.03  0.0421637
## 7  1e+02  0.03  0.0421637
## 8  1e+03  0.03  0.0421637
## 9  1e+04  0.03  0.0421637
## 10 1e+05  0.03  0.0421637
```

The optimal cost parameter is 0.001.

## SVM on Linear inseparable dataset

### Generate Linear inseparable dataset

For the second part of our experiements, we will generate a linear inseparable dataset, simiarly of 100 observation, with 2 feature, split 50/50. The negative class will be split into 2 clusters, each with a mean



higher or lower than the positive class.

```
# Set seed
set.seed(300)

# half and half with different means
positive_1 <- matrix(rnorm((data_points/2), mean = pos_mean), nrow = 50, ncol = 2)

positive_2 <- matrix(rnorm((data_points/2), mean = neg_mean), nrow = 50, ncol = 2)

set.seed(300)

negative_1 = matrix(c(rnorm((data_points/2), mean = neg_mean) + 3,
                      rnorm((data_points/2), mean = neg_mean)), nrow = 50, ncol = 2)

negative_2 = matrix(c(rnorm((data_points/2), mean = pos_mean) - 3,
                      rnorm((data_points/2), mean = pos_mean)), nrow = 50, ncol = 2)
```

We will again create simulated labels for our new dataset.

```
new_labels <- c(rep(1, 100), rep(-1, 100))
```


And again, we will bind and populate a dataframe.

```
new_obs <- rbind(positive_1, positive_2, negative_1, negative_2)

new_data <- data.frame(x = new_obs, y = as.factor(new_labels))
```

We visualise our dataset.

```
plot(new_obs, col=ifelse(new_labels > 0, 1, 2))
legend("topleft", c("Positive", "Negative"), col=seq(2), pch=1, text.col=seq(2))
```



Finally we can again split our data into training and test sets by a 70/30 split.

```
set.seed(300)

new_train_indices <- sample(200, 200 * 0.7)

new_training_set <- new_data[new_train_indices, ]

new_test_obs <- obs[-new_train_indices, ]
new_test_labels <- labels[-new_train_indices]
```

## Part 8: SVM model with a radial kernel, $\gamma = 1$ and $\text{cost} = 1$ .

To deal with our linear inseparable dataset, we will build a radial model, with a cost of 1 and gamma parameter also of 1.

```
non_linear_fit_1 = svm(y~., data = new_training_set, kernel = "radial", gamma = 1, cost = 1)
```

### a: Summary

```
summary(non_linear_fit_1)
```


```
##
## Call:
```

```
## svm(formula = y ~ ., data = new_training_set, kernel = "radial",
##      gamma = 1, cost = 1)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost: 1
##      gamma: 1
##
## Number of Support Vectors: 23
##
## ( 11 12 )
##
## Number of Classes: 2
##
## Levels:
## -1 1
```

The model has 28 support vectors.

b: Plot

```
plot(non_linear_fit_1, new_training_set)
```



### C: Test error

```
# predictions
prediction_non_linear_1 = predict(non_linear_fit_1, newdata = new_test_obs)

# error calculation
mean(prediction_non_linear_1 != new_test_labels)
```

```
## [1] 0.5
```

The model has and error of 0.5.

## Part 9: SVM model with a radial kernel, $\gamma = 1$ and $\text{cost} = 1e5$ .

We will follow the same process as part 8, build instead use a model with a cost of  $1e5$  and a  $\gamma$  value of  $-1$

```
# build model
non_linear_fit_1e5 = svm(y~., data = new_training_set, kernel = "radial", gamma = 1, cost = 1e5)
```

### a: Summary

```
summary(non_linear_fit_1e5)


##
## Call:
## svm(formula = y ~ ., data = new_training_set, kernel = "radial",
##      gamma = 1, cost = 1e+05)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:  1e+05
##      gamma:  1
##
## Number of Support Vectors:  9
##
##  ( 6 3 )
##
##
## Number of Classes:  2
##
## Levels:
##  -1  1
```

The model has used 12 support vectors.

### b: Plot

```
plot(non_linear_fit_1e5, new_training_set)
```

## SVM classification plot



C: Test error

```
# predictions
prediction_non_linear_1e5 = predict(non_linear_fit_1e5, newdata = new_test_obs)

# error calculation
mean(prediction_non_linear_1e5 != new_test_labels)

## [1] 0.5833333
```

The model gives an error of 0.4833333 - a slight improvement.

## Part 10: Optimal values for cost and gamma

We can use the tune function to find the optimal values for both parameters.

```
set.seed(1)

costs <- c(0.001, 0.01, 0.1, 1, 5, 10, 100, 1000, 10000, 1e5)

gammas <- c(0.5, 1, 1.5, 2, 2.5, 3, 3.5, 4)

tuned_radial <- tune(svm, y~., data = new_training_set, kernel = "radial", ranges = list(cost = costs, gamma = gammas))
```

Summarise the results.

```
summary(tuned_radial)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   0.1    0.5
##
## - best performance: 0.007142857
##
## - Detailed performance results:
##   cost gamma      error dispersion
## 1  1e-03    0.5 0.571428571 0.19047619
## 2  1e-02    0.5 0.571428571 0.19047619
## 3  1e-01    0.5 0.007142857 0.02258770
## 4  1e+00    0.5 0.007142857 0.02258770
## 5  5e+00    0.5 0.007142857 0.02258770
## 6  1e+01    0.5 0.014285714 0.03011693
## 7  1e+02    0.5 0.014285714 0.03011693
## 8  1e+03    0.5 0.014285714 0.03011693
## 9  1e+04    0.5 0.064285714 0.06254250
## 10 1e+05    0.5 0.064285714 0.06254250
## 11 1e-03    1.0 0.571428571 0.19047619
## 12 1e-02    1.0 0.571428571 0.19047619
## 13 1e-01    1.0 0.007142857 0.02258770
## 14 1e+00    1.0 0.007142857 0.02258770
## 15 5e+00    1.0 0.007142857 0.02258770
## 16 1e+01    1.0 0.014285714 0.03011693
## 17 1e+02    1.0 0.014285714 0.03011693
## 18 1e+03    1.0 0.050000000 0.05880519
## 19 1e+04    1.0 0.064285714 0.06254250
## 20 1e+05    1.0 0.064285714 0.06254250
## 21 1e-03    1.5 0.571428571 0.19047619
## 22 1e-02    1.5 0.571428571 0.19047619
## 23 1e-01    1.5 0.007142857 0.02258770
## 24 1e+00    1.5 0.007142857 0.02258770
## 25 5e+00    1.5 0.007142857 0.02258770
## 26 1e+01    1.5 0.014285714 0.03011693
## 27 1e+02    1.5 0.014285714 0.03011693
## 28 1e+03    1.5 0.050000000 0.05880519
## 29 1e+04    1.5 0.064285714 0.05270463
## 30 1e+05    1.5 0.064285714 0.05270463
## 31 1e-03    2.0 0.571428571 0.19047619
## 32 1e-02    2.0 0.571428571 0.19047619
## 33 1e-01    2.0 0.007142857 0.02258770
## 34 1e+00    2.0 0.007142857 0.02258770
## 35 5e+00    2.0 0.014285714 0.03011693
## 36 1e+01    2.0 0.014285714 0.03011693
## 37 1e+02    2.0 0.014285714 0.03011693
## 38 1e+03    2.0 0.057142857 0.05634362
```

```
## 39 1e+04 2.0 0.071428571 0.05832118
## 40 1e+05 2.0 0.071428571 0.05832118
## 41 1e-03 2.5 0.571428571 0.19047619
## 42 1e-02 2.5 0.571428571 0.19047619
## 43 1e-01 2.5 0.007142857 0.02258770
## 44 1e+00 2.5 0.007142857 0.02258770
## 45 5e+00 2.5 0.014285714 0.03011693
## 46 1e+01 2.5 0.014285714 0.03011693
## 47 1e+02 2.5 0.021428571 0.03450328
## 48 1e+03 2.5 0.057142857 0.05634362
## 49 1e+04 2.5 0.057142857 0.05634362
## 50 1e+05 2.5 0.057142857 0.05634362
## 51 1e-03 3.0 0.571428571 0.19047619
## 52 1e-02 3.0 0.571428571 0.19047619
## 53 1e-01 3.0 0.007142857 0.02258770
## 54 1e+00 3.0 0.007142857 0.02258770
## 55 5e+00 3.0 0.014285714 0.03011693
## 56 1e+01 3.0 0.014285714 0.03011693
## 57 1e+02 3.0 0.028571429 0.04994328
## 58 1e+03 3.0 0.050000000 0.05880519
## 59 1e+04 3.0 0.050000000 0.05880519
## 60 1e+05 3.0 0.050000000 0.05880519
## 61 1e-03 3.5 0.571428571 0.19047619
## 62 1e-02 3.5 0.571428571 0.19047619
## 63 1e-01 3.5 0.007142857 0.02258770
## 64 1e+00 3.5 0.007142857 0.02258770
## 65 5e+00 3.5 0.014285714 0.03011693
## 66 1e+01 3.5 0.014285714 0.03011693
## 67 1e+02 3.5 0.028571429 0.04994328
## 68 1e+03 3.5 0.042857143 0.06023386
## 69 1e+04 3.5 0.042857143 0.06023386
## 70 1e+05 3.5 0.042857143 0.06023386
## 71 1e-03 4.0 0.571428571 0.19047619
## 72 1e-02 4.0 0.571428571 0.19047619
## 73 1e-01 4.0 0.007142857 0.02258770
## 74 1e+00 4.0 0.007142857 0.02258770
## 75 5e+00 4.0 0.014285714 0.03011693
## 76 1e+01 4.0 0.014285714 0.03011693
## 77 1e+02 4.0 0.028571429 0.04994328
## 78 1e+03 4.0 0.028571429 0.04994328
## 79 1e+04 4.0 0.028571429 0.04994328
## 80 1e+05 4.0 0.028571429 0.04994328
```

We can see that the best value for cost is 0.1 and the best value for gamma is 0.5.

## Part 11: Test error on optimal model

Finally, we can check the test error on our optimal model.

```
prediction_optimal_radial = predict(tuned_radial$best.model, newdata = new_test_obs)

mean(prediction_optimal_radial != new_test_labels)
```

```
## [1] 0.5
```